



SQL AND ORACLE
SIT103 Lecture 6

1

CORRELATED SUB-QUERY

- This is also a nested query with one difference. The nested query refers to a field in an outer query.
- Since a reference is made to an outer query field, the query can NOT be executed once before outer query.

CORRELATED SUB-QUERY - EXAMPLE

- The inner query refers to a field from the outer query.

```
SELECT STUDENT_NO, SURNAME, GIVEN
FROM STUDENT
WHERE 3 =
  (SELECT COUNT(*)
   FROM STUD_COURSE
   WHERE STUD_COURSE.STUDENT_NO
     = STUDENT.STUDENT_NO);
```

Same Table

CORRELATED SUB-QUERY - EXECUTION

- Execution of correlated queries:
 - For each record in the outer query:
 - execute inner sub-query
 - produce result
 - use result in the execution of the outer query

EXISTS AND SUB-QUERIES

- Check if the sub-query returns any answer.
- NOT version also possible

```
SELECT STUDENT_NO, SURNAME, GIVEN
FROM STUDENT
WHERE EXISTS
  (SELECT *
   FROM STUD_COURSE
   WHERE STUD_COURSE.STUDENT_NO
     = STUDENT.STUDENT_NO);
```

NOTE: Can also use EXISTS, NOT EXISTS, IN, NOT IN instead of "<".

PL/SQL (106)

- PL/SQL is procedural programming language roughly based on SQL.
- It has similar constructs to any other programming language such as variables, IF statement and Loops
- It also has special constructs such as cursors to allow looping through a table one row at a time.
- Results of SQL statements such as SELECT are not displayed to the user, but instead put into variables.

EXAMPLE OF PL/SQL

```
DECLARE
-- variable declaration
message varchar2(20):= 'Hello, World!';
BEGIN
/* * PL/SQL executable statement(s) */
dbms_output.put_line(message);
END;
```

STORED PROCEDURES

- A Stored procedure is a named block of procedural code which is compiled and stored on the server, in the schema of the user who created it.
- It is the same, conceptually, as a subroutine in any other programming language - a method in Java.
- It can be passed parameters.
- It can be called with the EXEC command from another procedure or another program.
- Stored procedures have no user input or output - they would normally store results into tables.

STORED PROCEDURE EXAMPLE

```
CREATE OR REPLACE
PROCEDURE addNewInvoice (givenCustomer
NUMBER) AS
BEGIN
INSERT INTO INVOICE VALUES(
INVNUM.NEXTVAL, SYSDATE, givenCustomer);
END ;
/
```

- To call this procedure from the SQLPlus prompt, from another procedure, or from VB or Java:
EXEC addNewInvoice(127) ;

STORED PROCEDURE NOTES

- Note the parameter in the previous example.
- You can type the code directly at the SQL-Plus prompt, or into a file, as with a SQL query.
- You signal the end of a PL/SQL block with a full stop ., or slash / on a line by itself.
- When you start the file, Oracle attempts to compile and store your code block.



USER DEFINED FUNCTIONS

- Functions are similar to Stored Procedures.
- They are named and stored on the server in the schema of the person who created them
- They can be made available to other users.
- Syntax rules are exactly the same
- Main difference a function returns a value.



USER DEFINED FUNCTION EXAMPLE

```
CREATE OR REPLACE
FUNCTION theYear( theDate DATE) RETURN
NUMBER IS
BEGIN
    RETURN TO_NUMBER( TO_CHAR( theDate, 'YYYY'
    ));
END theYear;
/
```



FUNCTION USE

- Users with execute privilege on function theYear could then use it in a SQL statement.

```
SELECT theYear( DOB) FROM STUDENT ;  
SELECT theYear(SYSDATE) + 20 FROM DUAL  
;
```
- The user who created the function automatically has all privileges on it.
- S/he can grant the execute privilege to others:

```
GRANT EXECUTE ON theYear TO FRED;
```



TRIGGERS

- Triggers are similar to Stored Procedures and Functions in that :
 - they are named blocks of PL/SQL code
 - they are compiled and stored on the database server
- What differs - they are not called explicitly by a user, procedure, function or program.
- When triggers are defined, they are attached to a particular table, for particular events such as INSERT, UPDATE or DELETE.



TRIGGERS

- Triggers are fired when the corresponding triggering event happens on the table.
 - Eg: a user issues an INSERT command, or UPDATE or DELETE
- They are also often used to implement complex auditing and updating:
 - Eg: to update a stock on hand value when a sale occurs



TRIGGER SYNTAX

- The general syntax is :

```
CREATE OR REPLACE TRIGGER trigger_name
{BEFORE | AFTER} triggering_event ON
table_reference
[FOR EACH ROW [WHEN trigger_condition] ]
trigger_body ;
```

- You can query the data dictionary view user_triggers to see trigger details.

```
SELECT trigger_type, table_name,
triggering_event
FROM USER_TRIGGERS ;
```

TRIGGER EXAMPLE 1 (NEW)

```
CREATE OR REPLACE TRIGGER atleast15
AFTER INSERT OR UPDATE OF DOB ON
STUDENT
FOR EACH ROW
DECLARE thedob DATE ;
BEGIN
thedob := :new.DOB;
IF (( SYSDATE - thedob ) < ( 15 * 365.25)) THEN
RAISE_APPLICATION_ERROR( -20000,
'Students must be at least 15 years old');
END IF;
END ;
/
```

:NEW AND :OLD (NEW)

- **:new** and **:old** are pseudo tables with exactly the same structure as the trigger table.
- These pseudo tables hold the replacing (after) and replaced (before) values.
- You can reference them only in a trigger.
- Each RDBMS which supports triggers provides similar pseudo tables. MS SQL Server calls them inserted and deleted.

TRIGGER EXAMPLE 2

- Suppose we want to record each time a student changes programmes. We log these changes into a log table created thus:

```
CREATE TABLE PGM_CHANGE  
( STUDENT_NO      CHAR(8),  
  FROM_PGM        VARCHAR2(6),  
  TO_PGM           VARCHAR2(6),  
  CHANGE_DATE      DATE );
```

TRIGGER EXAMPLE 2

```
CREATE OR REPLACE TRIGGER logPgmChange  
BEFORE UPDATE ON STUDENT  
FOR EACH ROW  
WHEN ( new.PGMCODE != old.PGMCODE )  
  -- don't use colon in WHEN clause, in trigger body  
  only  
BEGIN  
  INSERT INTO PGM_CHANGE VALUES (  
    :old.STUDENT_NO, :old.PGMCODE,  
    :new.PGMCODE, SYSDATE);  
END logPgmChange;
```

INDEXES

- An index speeds up searching, sorting and joining operations
- Indexes slow down updates, however.
- Index is simply created.
- System handles all updates to the index.
- The system decides if the index will be used.
- A command cannot specify the use of an index.

INDEXES

- CREATE INDEX STUDDSURN
ON STUDENT(SURNAME);
- CREATE UNIQUE INDEX STUCRSPRIMARY
ON
STUD_COURSE(STUDENT_NO,COURSE_CODE);

Note: above creates a primary key type index – sometimes used for candidate keys



INDEXES

- Create Indexes on columns used for:
 - Primary keys (unique)
 - Foreign keys
 - search fields
 - ordering fields
- Do NOT create indexes on:
 - Fields with few different values
 - Small tables
 - NULL values



ORACLE SEQUENCES

- Sequences are used to generate unique numbers for primary keys. Oracle does not have a type to do this.

```
CREATE SEQUENCE TRANSIDSEQ  
START WITH 1000 INCREMENT BY 1;
```

- The sequence is used explicitly within an INSERT statement to get the next value by referencing the NEXTVAL variable.



ORACLE SEQUENCES

INSERT INTO TRANSACTION
VALUES(TRANSIDSEQ.NEXTVAL,.....)

- Each reference to TRANSIDSEQ.NEXTVAL increments it.
- To see the current value without incrementing it, you can reference TRANSIDSEQ.CURRVAL.



SYSTEM VIEWS (NEW)

- To find out about various objects in the database, there are hundreds of views defined that display system information.
- Many system views are only usable by the DBA but there are many that a normal user can query to find out about their objects such as tables, views, indexes etc.
- It is useful to know some of these.



SYSTEM VIEWS - SOME EXAMPLES (N)

- USER_TABLES
- USER_VIEWS
- USER_CONSTRAINTS
- USER_INDEXES
- USER_SYNONYMS
- USER_TRIGGERS



DUAL SYSTEM TABLE (NEW)

- **DUAL** is a table with one column, DUMMY VARCHAR2(1), and one row with the value 'X'.
- It is useful for computing a single constant expression with the SELECT command.
- Because DUAL has only one row, the constant is only returned once. Eg:

SELECT SYSDATE FROM DUAL ;

- You could select SYSDATE from any table, but Oracle would return SYSDATE once for each row.


