



# **DATABASE ADMINISTRATION**

**SIT103 Lecture 10**

**Transaction Processing, Locking and  
Deadlocks, Backups and Recovery,  
Performance Issues**

# DATA ADMINISTRATION

- Manage the data needs of the entire organization so all data is available.
- Database Planning
- Data Analysis, Design and Implementation
- Data Maintenance and Protection
- Education and Training
- Support to end users
- High level management role



# DATABASE ADMINISTRATION

- Manage physical database servers.
- More of a technical, hands on role.
- Installation and configuration of DBMS
- Data Analysis, Design and Implementation
- Performance Tuning
- Managing Security and Policies
- Managing backups and recovery

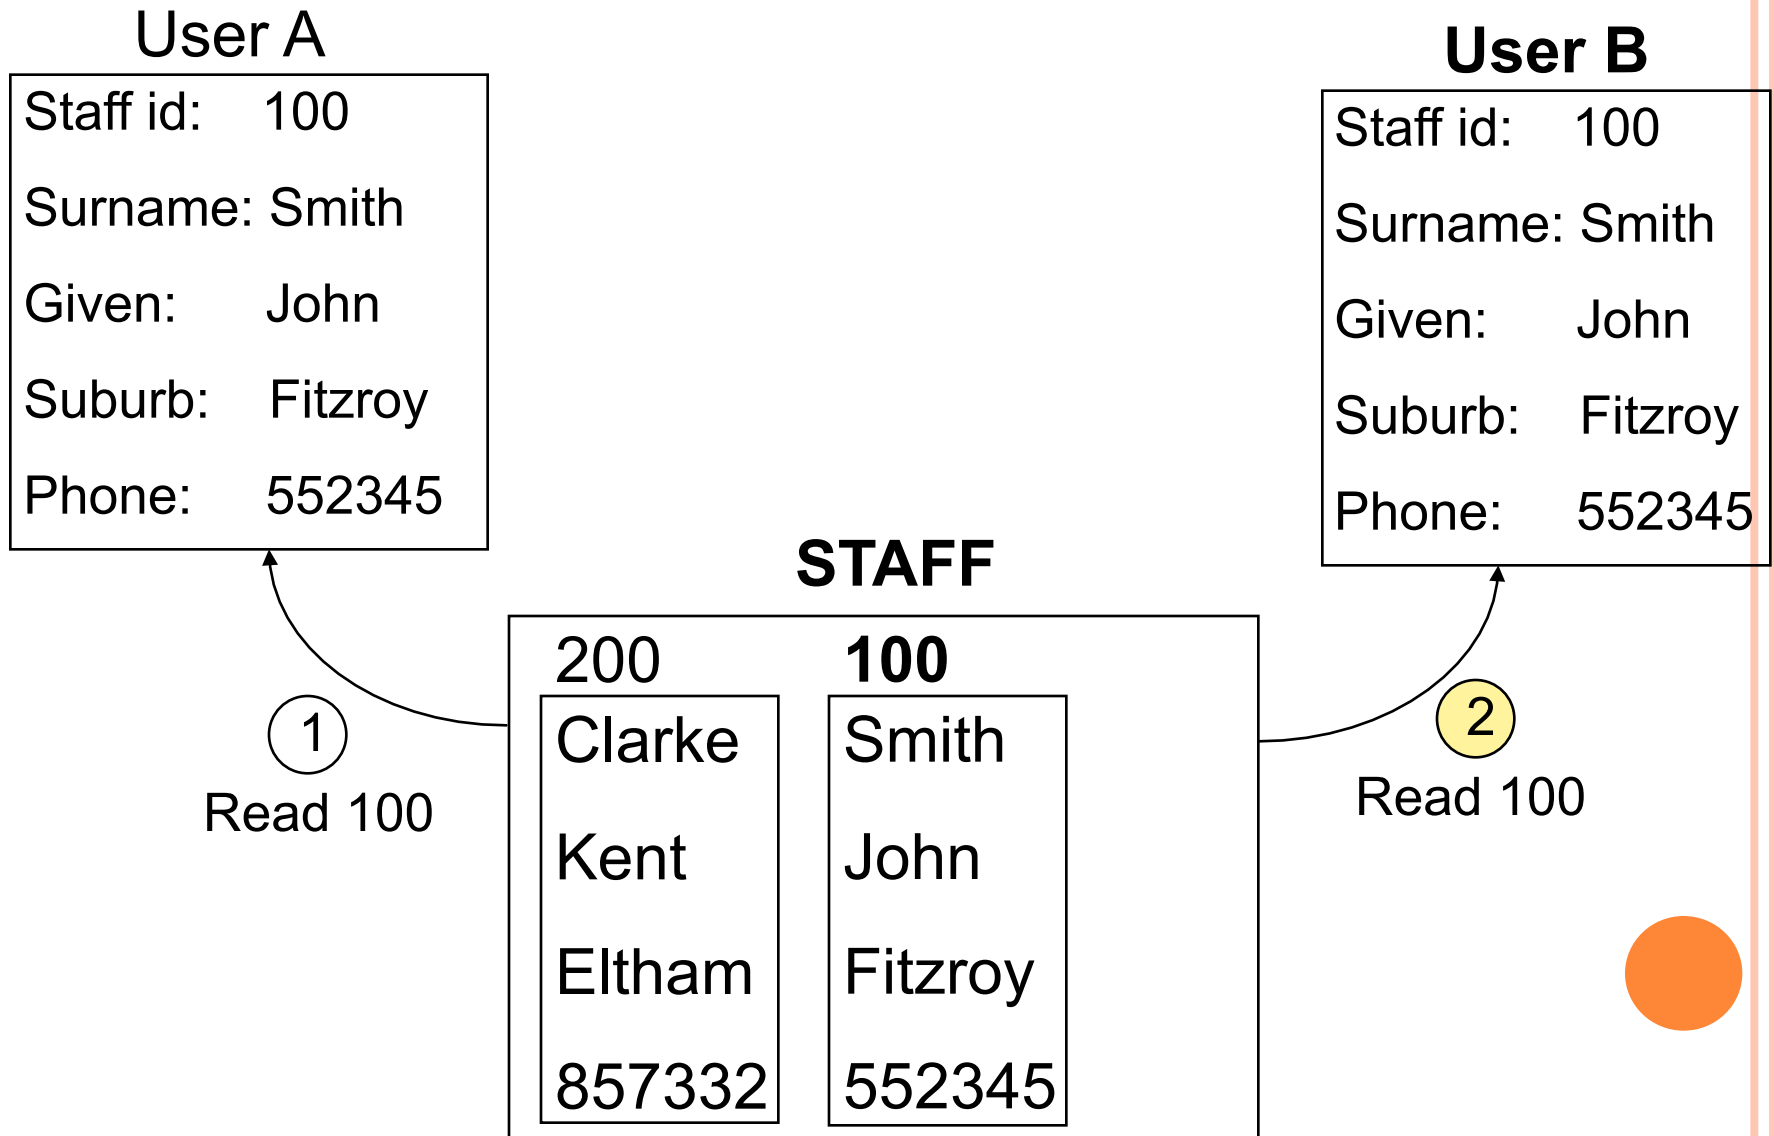


# CHANGING ROLES

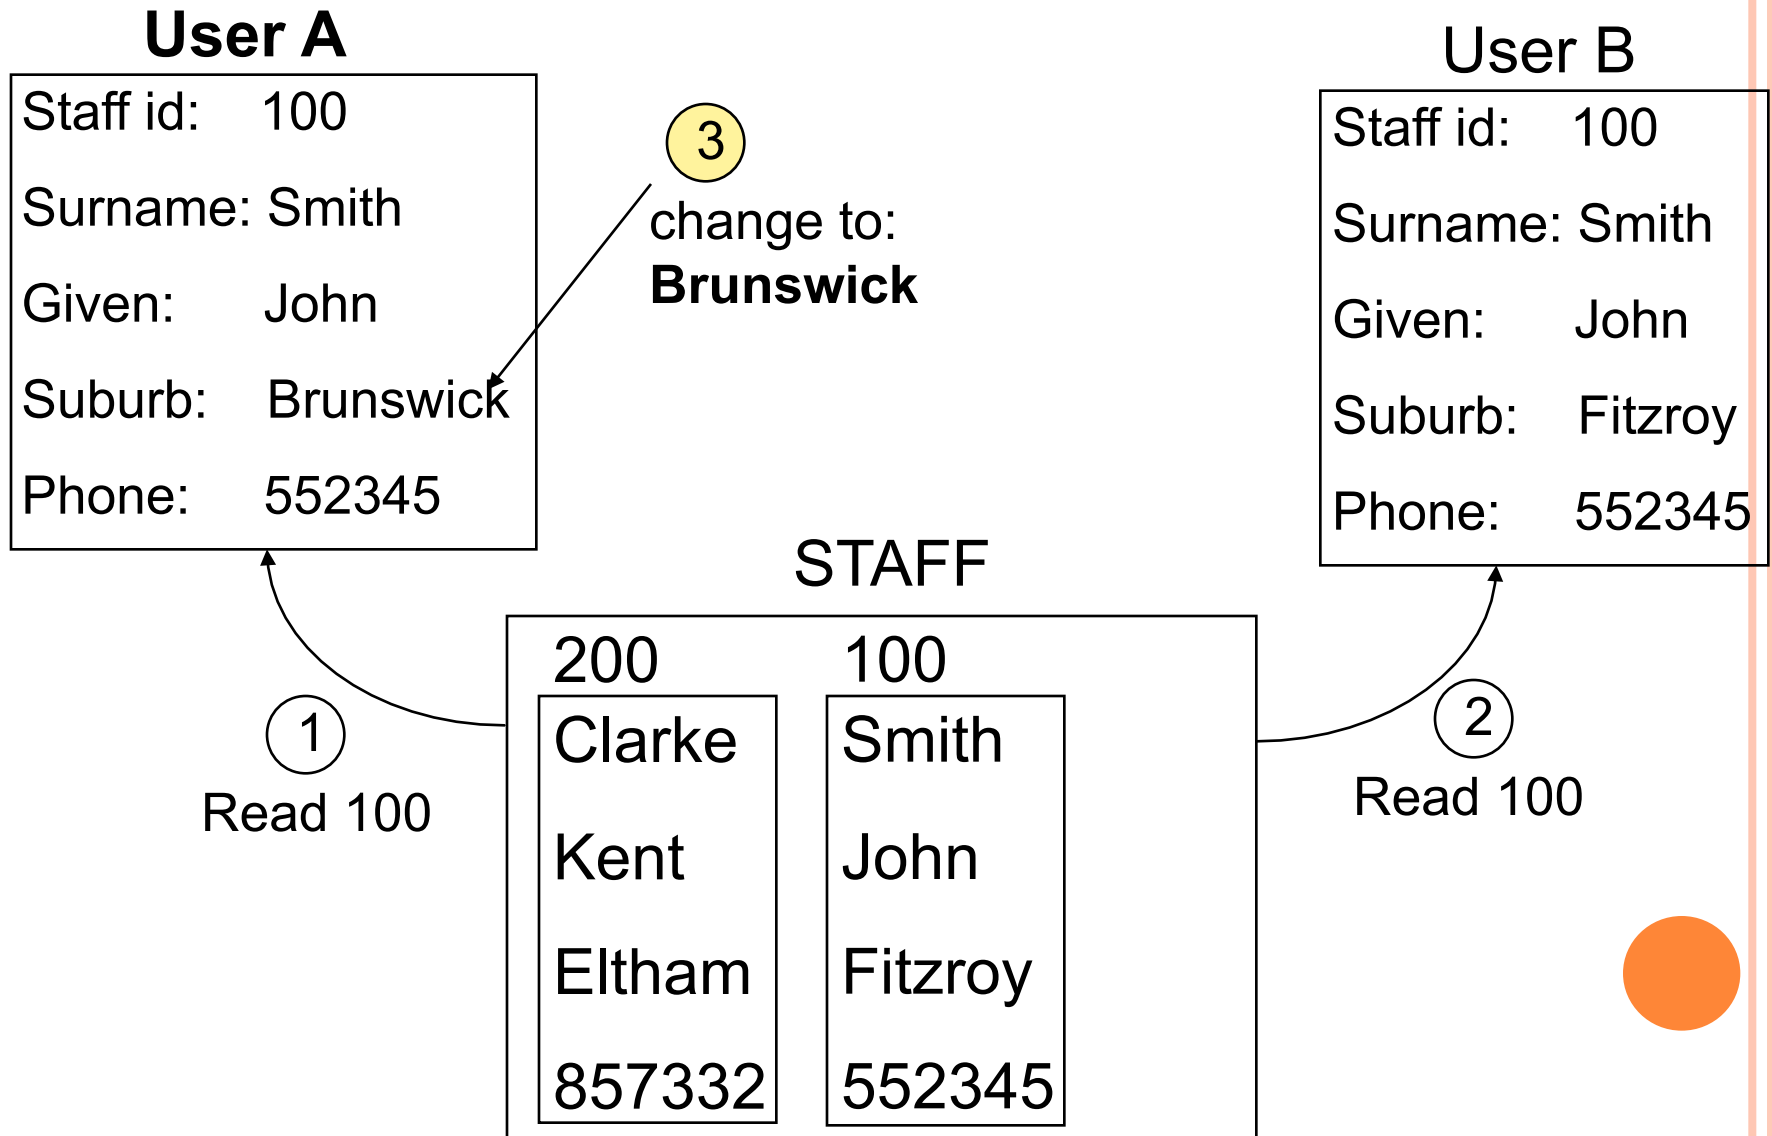
- In larger companies, these two roles may be performed by distinct people or groups.
- In smaller organisations or departments, one person will often perform both roles.
- DBA roles are becoming specialized:
  - Product specific (eg: Oracle DBA)
  - Overlap with developers (server programming)
  - Database types (eg: data warehouses)
  - Specific App Software (eg: SAP or PeopleSoft)



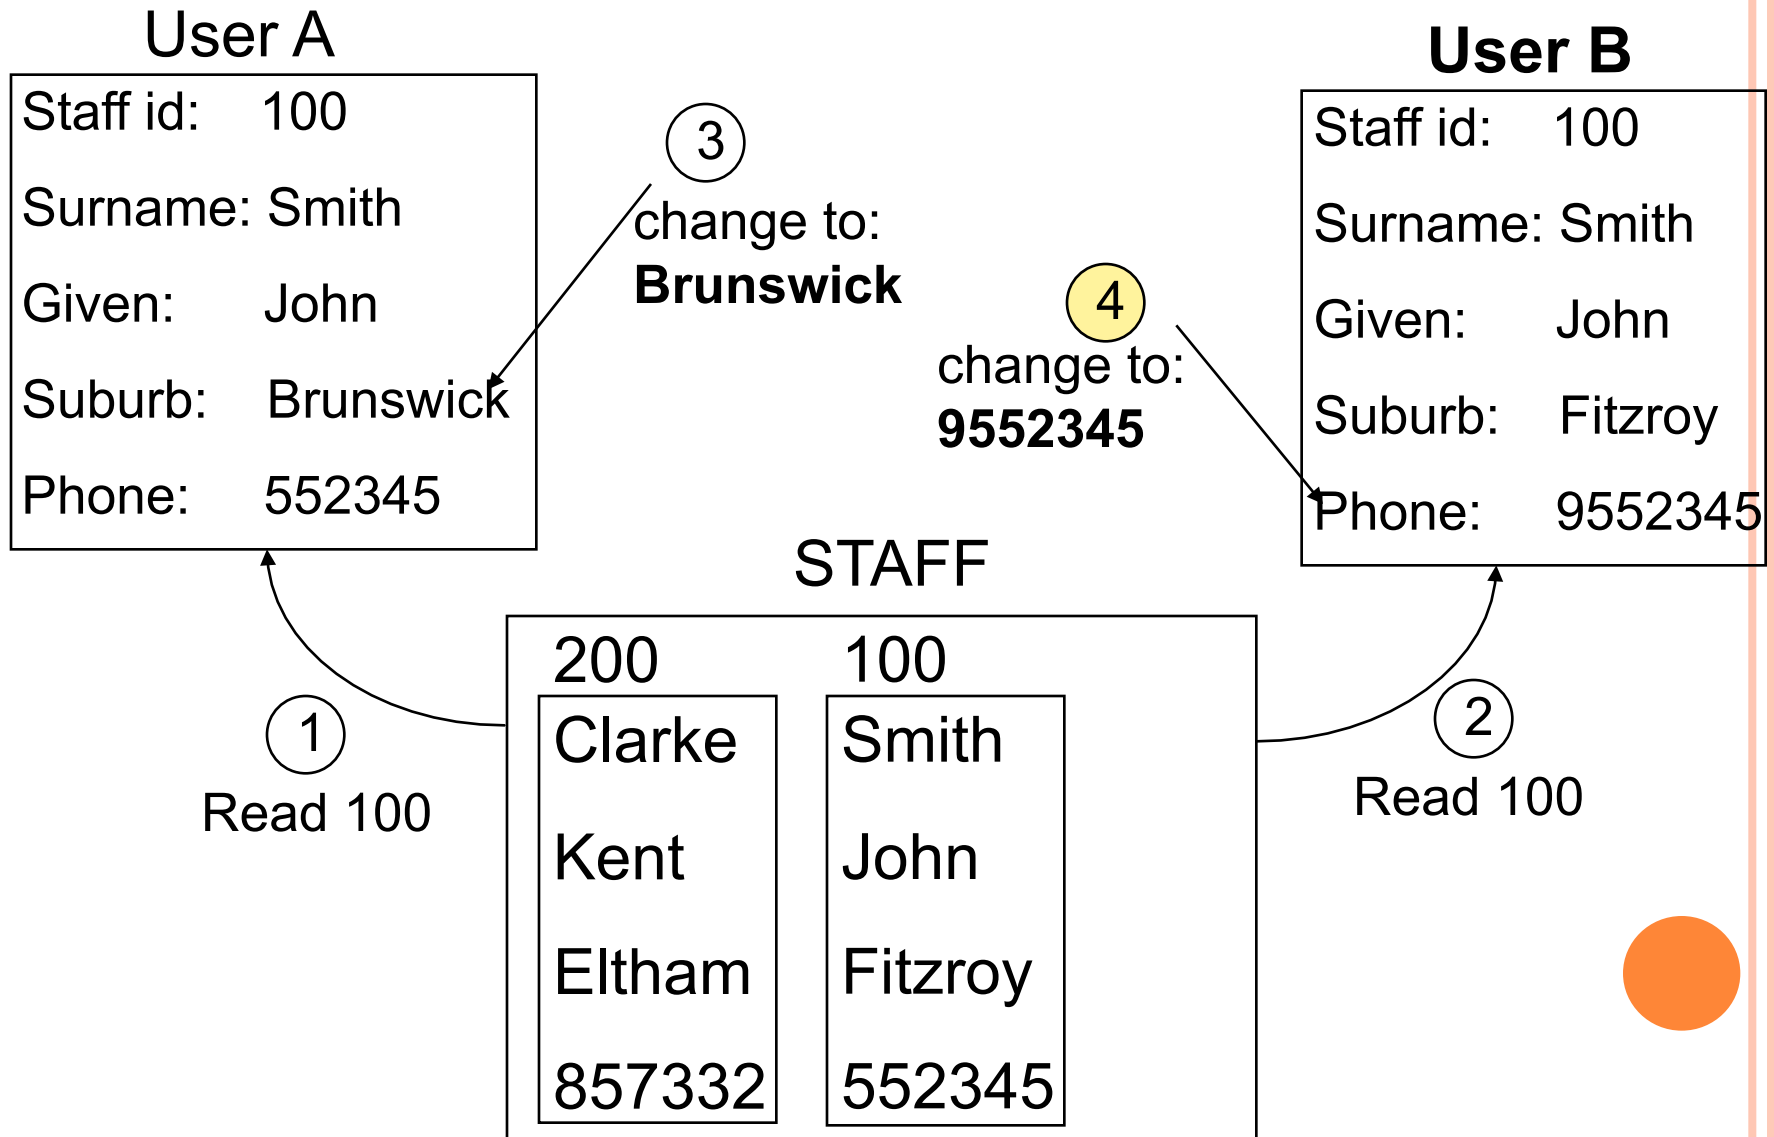
# Transaction Processing - no locking



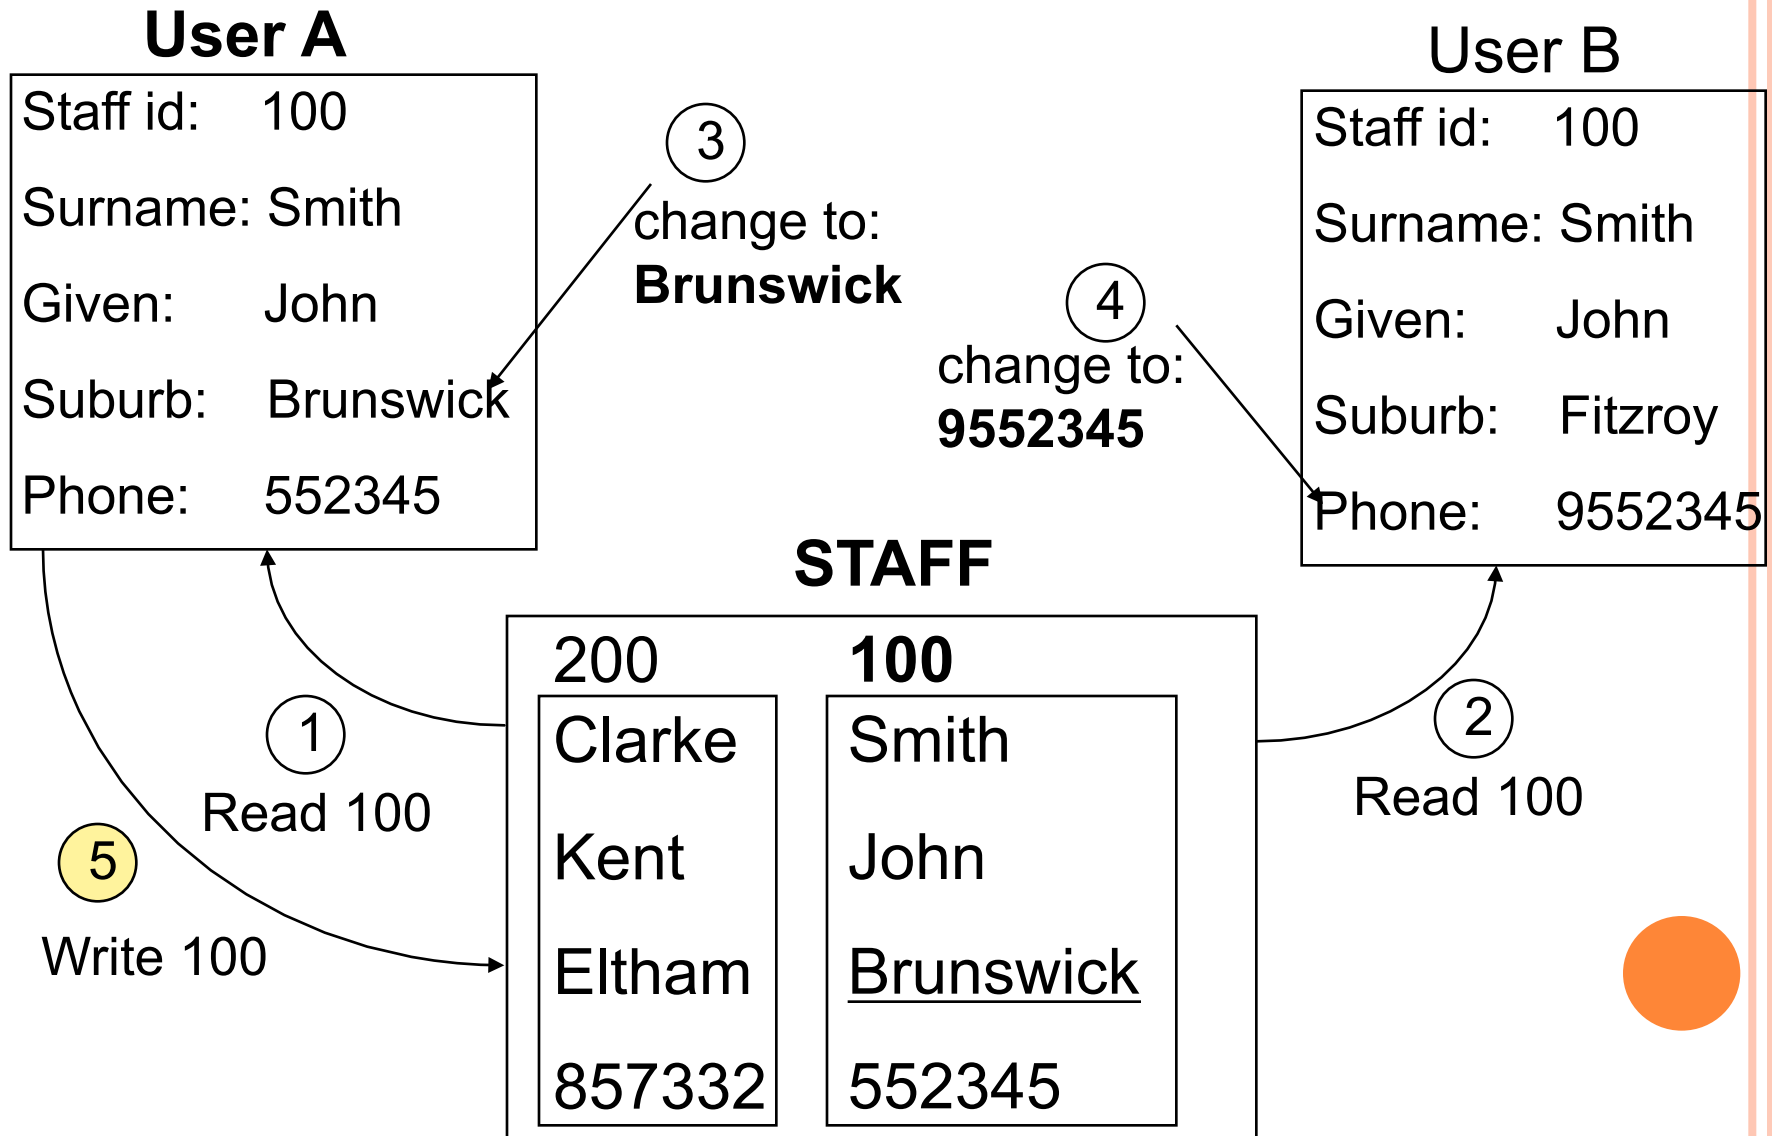
# Transaction Processing - no locking



# Transaction Processing - no locking

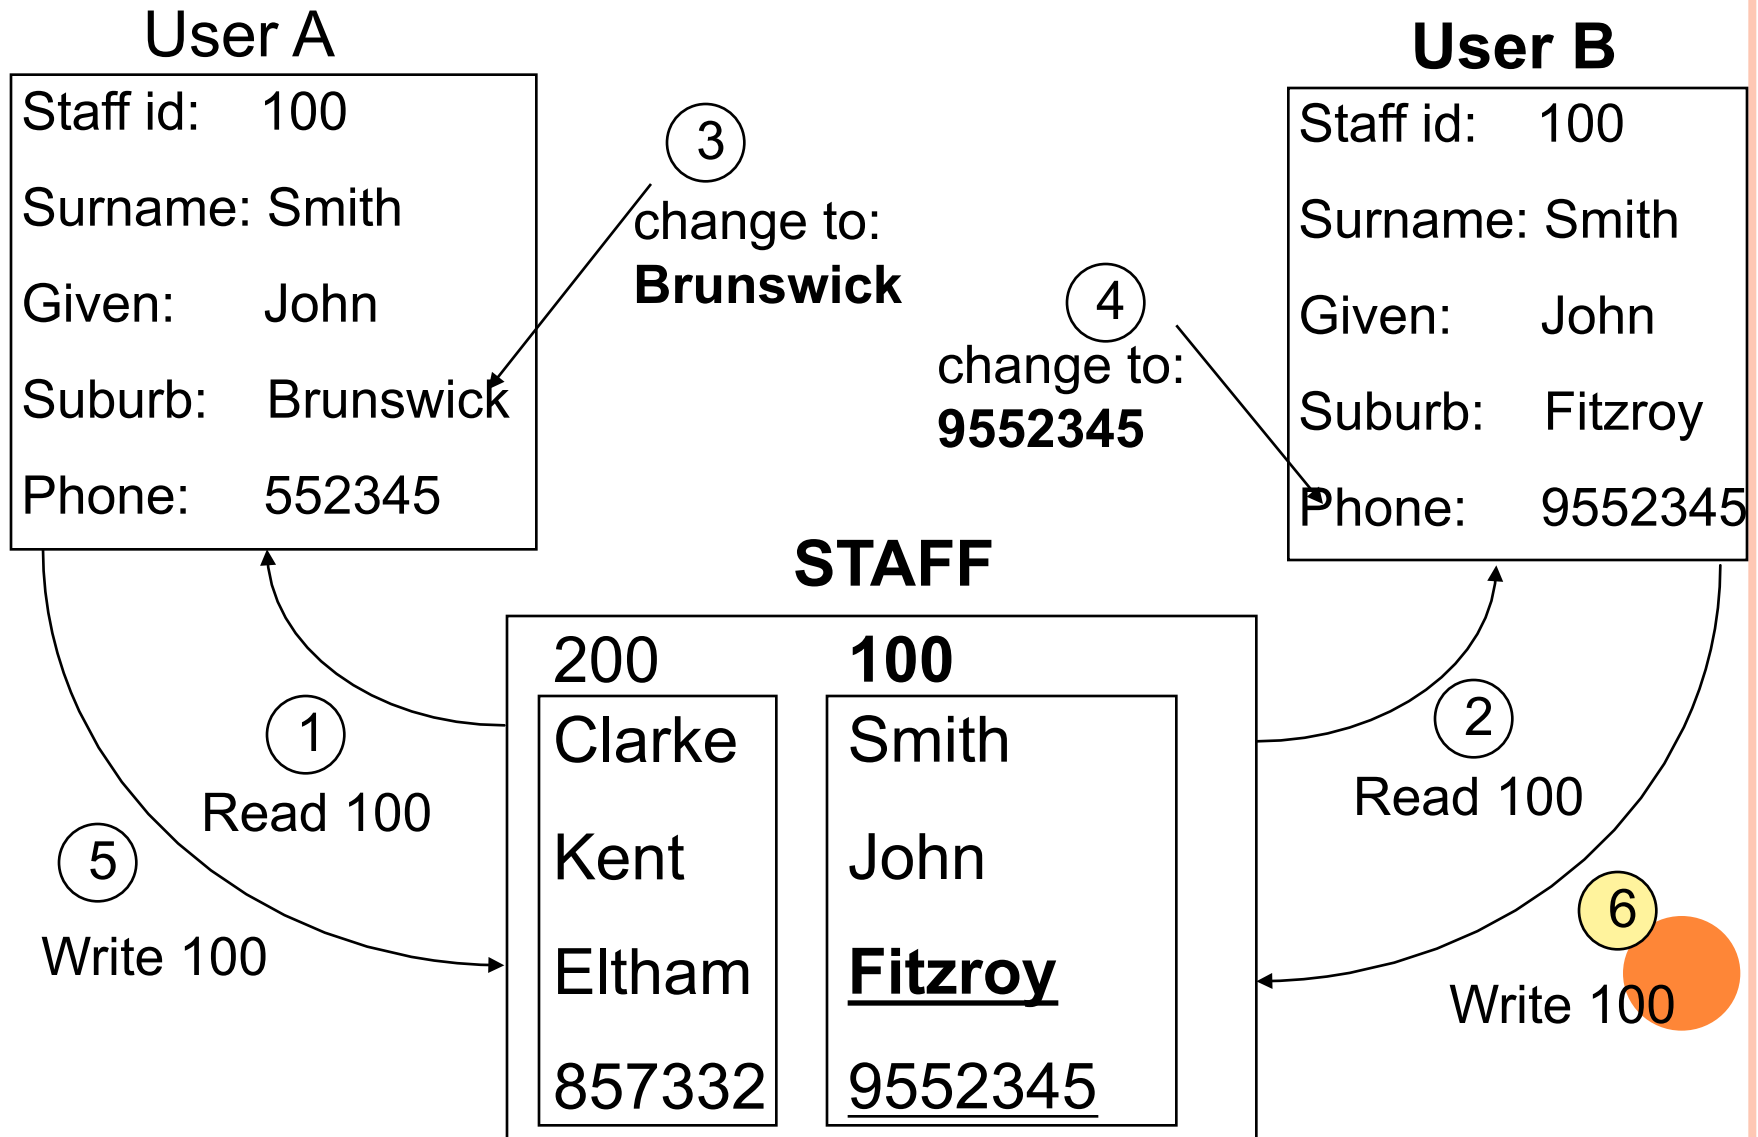


# Transaction Processing - no locking





# “known as the LOST UPDATE PROBLEM”



# LOCKING

- How do we solve this problem?
  - When a user/process wants to use a resource, LOCK IT.
  - No other user/process may change it until it is UNLOCKED.
- if used LOCKING in previous example order of execution would change, lets us see what happens.....



# Transaction Processing - with locking

## User A

Staff id: 100

Surname: Smith

Given: John

Suburb: Brunswick

Phone: 552345

2

change to:  
**Brunswick**

## User B

Staff id:

Surname:

Given:

Suburb:

Phone:

## STAFF

200

Clarke

Kent

Eltham

857332

100

Smith

John

Brunswick

552345

1

Read 100

3

Write 100



# Transaction Processing - with locking

## User A

Staff id: 100  
Surname: Smith  
Given: John  
Suburb: Brunswick  
Phone: 552345

## User B

Staff id: 100  
Surname: Smith  
Given: John  
Suburb: Brunswick  
Phone: 9552345

5  
change to:  
**9552345**

## STAFF

200	100
Clarke	Smith
Kent	John
Eltham	Brunswick
857332	<u>9552345</u>

4  
Read 100

6  
Write 100

# LEVEL OF LOCKING

1. **Database** - backup
2. **Table** - batch work (updates)
3. **Block/Page** - most common
4. **Record** - 1 record, most commonly required



# LOCK TYPES

## ○ Shared Locks

- other users/processes can read but not update; i.e. allowing querying;

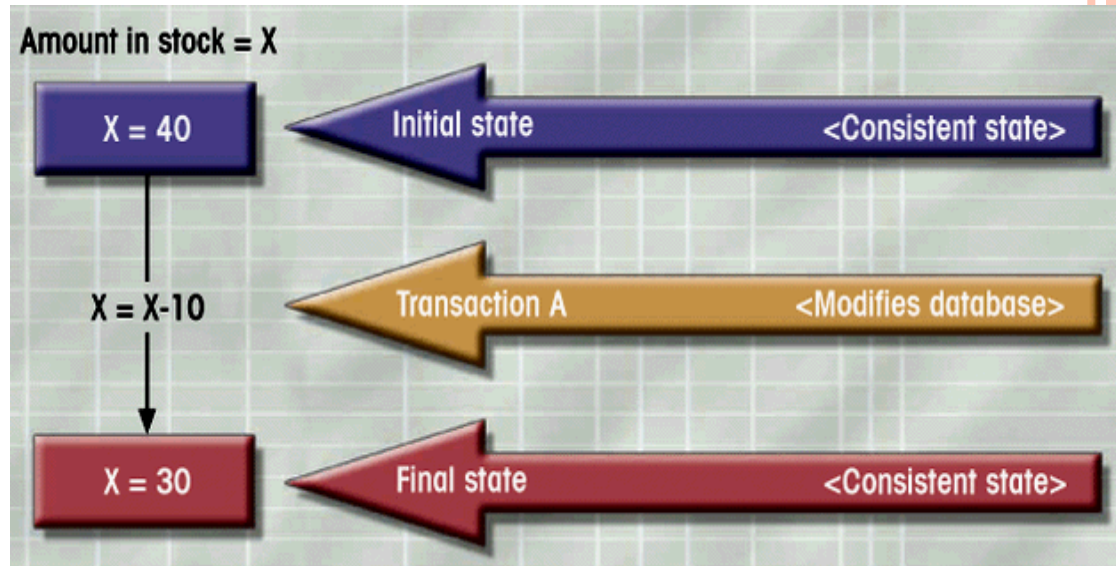
## ○ Exclusive Locks

- other users can do NOTHING.



# WHAT IS A TRANSACTION?

- Logical unit of work
  - Must be entirely completed or aborted
  - No intermediate states



# TRANSACTION

- Is a single business operation.
- A whole unit of work **MUST** be applied or **NONE** will be applied to database.
- eg. Transfer \$100 from AccA to AccB.

TRANSACTION

- Read AccA
- Subtract \$100 from AccA
- Write AccA
- Read AccB
- Add \$100 to AccB
- Write AccB





# DESIRABLE PROPERTIES OF TRANSACTIONS

## ○ Atomicity

- All transaction operations must be completed
- Incomplete transactions aborted

## ○ Consistency

- Maps one consistent database state to another

## ○ Isolation

- Transaction data cannot be reused until its execution complete

## ○ Durability

- Permanence of consistent DB state once a transaction commits



# EXAMPLE

- Assume the following SQL
  - Consistent state after transaction is executed
  - No changes made to database

```
SELECT ACC_NUM, ACC_BALANCE  
FROM CHECKACC  
WHERE ACC_NUM = '0908110638';
```



## EXAMPLE (CONT'D)

- Register credit sale of 100 units of product X to customer Y for \$500
  - State is consistent if both transactions complete
  - DBMS does not guarantee this

applied

```
UPDATE PRODUCT
SET PROD_QOH = PROD_QOH - 100
WHERE PROD_CODE = 'X';
```

```
UPDATE ACCT_RECEIVABLE
SET ACCT_BALANCE = ACCT_BALANCE + 500
WHERE ACCT_NUM = 'Y';
```

# TRANSACTION MANAGEMENT IN SQL

- Transaction support
  - COMMIT
  - ROLLBACK
- A transaction sequence must continue until:
  - COMMIT statement is reached
  - ROLLBACK statement is reached
  - End of a program reached
  - Program reaches abnormal termination



# DEADLOCKS

- Deadlocks arise from more than one process holding a resource and requesting other resources which are in turn held by other processes, for example:

## User A

transfer \$100 from X to Y

Read X {X is locked}

Subtract \$100 from X

Read Y

- Y is locked by user A, waiting

## User B

transfer \$100 from Y to X

Read Y {Y is locked}

Subtract \$100 from Y

Read X

- X is locked by user A, waiting

# DEADLOCK RESOLUTION

- Pessimistic Locking Scheme
  - every process must lock ALL resources required first.
  - if cannot, unlock everything and have another go later
  - performance slow due to overhead of locking everything.



# DEADLOCK RESOLUTION

## ◦ Deadlock Watchdog

- Let processes do what they want/please
- program keeping an eye on locks
- If DEADLOCK recognized then kill off processes until resolved.



# DEADLOCK RESOLUTION

- Optimistic Locking Scheme (versioning)
  - read and use records without concern
  - before writing at end of transaction
    - re-read record(s)
    - if not changed since 1st read the WRITE Ok
    - if has changed then abort whole transaction!
  - efficient - performance high; if over 90%, no change then QUICK1





# BACKUPS

- If your computer room has a fire and everything is lost, how long would it take to replace:
  - Hardware
  - Operating System
  - DBMS (eg: Oracle)
  - Data in the database
- Which is the most valuable ??



# DATABASE RECOVERY - BACKUP

## ○ Full Backup

- everything is backed up on some sort of magnetic media.

## ○ Partial Backup

- backup changes made since last FULL Backup was made.



# DATABASE RECOVERY - BACKUP


- Recommended backup philosophy
  - FULL Backup each week
  - PARTIAL backups each day
  - ROTATING the media used
  - OFFSITE media storage



# DATABASE RECOVERY - LOGS

- Journals/Transaction Logs
- Log enough information about a change to allow it to be applied again, if required.
  - before and after images of record changed; and/or
  - details about each transaction;
  - log start and finish of transaction.

100
102
104
102
100
101
100




# DATABASE RECOVERY - CHECKPOINTS

- Allows all transactions to finish, at a particular time.
- Write a checkpoint, as a log entry. This gives us a point at which no transaction are unfinished.
- Useful for recovery, if a problem occurs with the database.

100
102
104
102
100
101
100
Checkpoint
105


# DATABASE RECOVERY - ROLLBACK

- If database intact
  - can undo changes until checkpoint because OK point.
  - undo unfinished transaction back to a safe point (usually checkpoint)



# DATABASE RECOVERY - ROLLFORWARD

- database retrieved from backup
- re-apply transactions up to latest checkpoint.



# DATABASE - SECURITY

- more sensitive data being stored
- more people given access to database
- security is defined as protection of the database against accidental or intentional loss, destruction or misuse.





# DATABASE - SECURITY

- Data management software usually provides the following security features:
  - creation of views to restrict user access;
  - authorization rules;
  - programs further constraining or limiting database access;
  - encryption of data;
  - identify user attempting any database access (biometric, smartcard).

